Linear estimates and LS–means

Søren Højsgaard and Ulrich Halekoh

doBy version 4.6.8 as of 2020-11-10

Contents

1 Introduction

1.1 Linear functions of parameters

A linear function of a p–dimensional parameter vector β has the form

$$C = L\beta$$

where L is a $q \times p$ matrix which we call the *Linear Estimate Matrix* of simply *LE-matrix*. The corresponding linear estimate is $\hat{C} = L\hat{\beta}$. A linear hypothesis has the form $H_0 : L\beta = m$ for some q dimensional vector m.
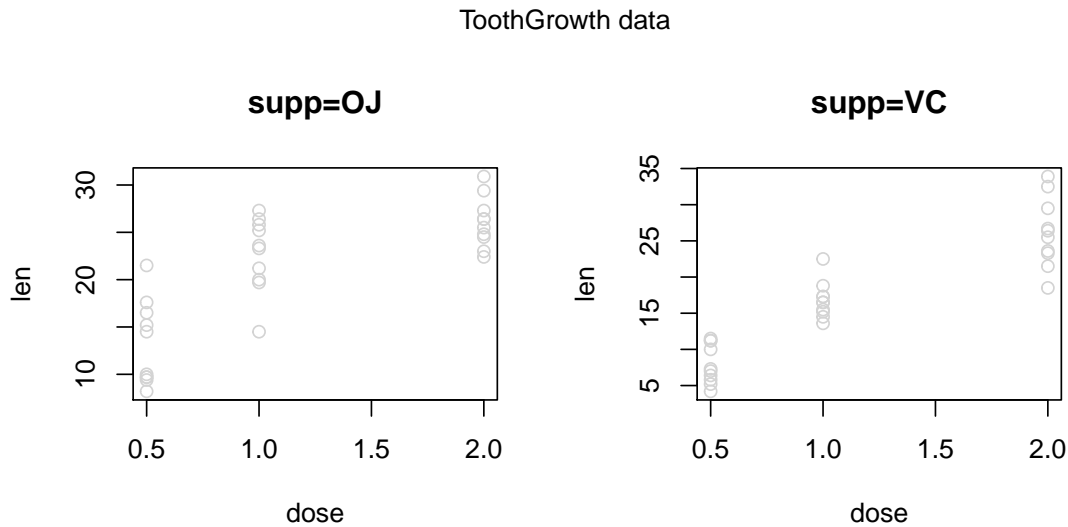
ToothGrowth data

supp=OJ   **supp=VC**

Figure 1: Plot of length against dose for difference sources of vitamin C.

1.2 Tooth growth

The response is the length of odontoblasts cells (cells responsible for tooth growth) in 60 guinea pigs. Each animal received one of three dose levels of vitamin C (0.5, 1, and 2 mg/day) by one of two delivery methods, (orange juice (coded as OJ) or ascorbic acid (a form of vitamin C and (coded as VC)).

```
head(ToothGrowth, 4)


##    len supp dose
## 1  4.2   VC  0.5
## 2 11.5   VC  0.5
## 3  7.3   VC  0.5
## 4  5.8   VC  0.5


ftable(xtabs(~ dose + supp, data=ToothGrowth))


##      supp OJ VC
## dose
## 0.5       10 10
## 1         10 10
## 2         10 10
```

The interaction plot suggests a mild interaction which is supported by a formal comparison:

```
ToothGrowth$dose <- factor(ToothGrowth$dose)
head(ToothGrowth)


##    len supp dose
## 1  4.2   VC  0.5
## 2 11.5   VC  0.5
```

Figure 2: Interaction plot between dose and source of vitamin C.

```
## 3   7.3   VC   0.5
## 4   5.8   VC   0.5
## 5   6.4   VC   0.5
## 6  10.0   VC   0.5


tooth1 <- lm(len ~ dose + supp, data=ToothGrowth)
tooth2 <- lm(len ~ dose * supp, data=ToothGrowth)
anova(tooth1, tooth2)


## Analysis of Variance Table
##
## Model 1: len ~ dose + supp
## Model 2: len ~ dose * supp
##   Res.Df RSS Df Sum of Sq    F Pr(>F)
## 1     56 820
## 2     54 712  2       108 4.11  0.022 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

2   Computing linear estimates

For now, we focus on the additive model:

```
tooth1


##
## Call:
## lm(formula = len ~ dose + supp, data = ToothGrowth)
##
```

```
## Coefficients:
## (Intercept)         dose1         dose2        suppVC
##       12.45          9.13         15.50         -3.70
```

Consider computing the estimated length for each dose of orange juice (OJ): One option: Construct the LE–matrix L directly:

```
L <- matrix(c(1, 0, 0, 0,
              1, 1, 0, 0,
              1, 0, 1, 0), nrow=3, byrow=T)
```

Then do:

```
c1 <- linest(tooth1, L)
c1
```

```
## Coefficients:
##      estimate std.error statistic     df p.value
## [1,]   12.455     0.988    12.603 56.000       0
## [2,]   21.585     0.988    21.841 56.000       0
## [3,]   27.950     0.988    28.281 56.000       0
```

We can do:

```
summary(c1)
```

```
## Coefficients:
##      estimate std.error statistic     df p.value
## [1,]   12.455     0.988    12.603 56.000       0
## [2,]   21.585     0.988    21.841 56.000       0
## [3,]   27.950     0.988    28.281 56.000       0
##
## Grid:
## NULL
##
## L:
##      [,1] [,2] [,3] [,4]
## [1,]    1    0    0    0
## [2,]    1    1    0    0
## [3,]    1    0    1    0
```

```
coef(c1)
```

```
##   estimate std.error statistic df   p.value
## 1    12.45    0.9883     12.60 56 5.490e-18
## 2    21.58    0.9883     21.84 56 4.461e-29
## 3    27.95    0.9883     28.28 56 7.627e-35
```

```
confint(c1)
```

```
##    0.025 0.975
## 1 10.48 14.43
## 2 19.61 23.56
## 3 25.97 29.93
```

3 Automatic generation of L

The matrix L can be generated as follows:

```
L <- LE_matrix(tooth1, effect="dose", at=list(supp="OJ"))
L
```

```
##      (Intercept) dose1 dose2 suppVC
## [1,]           1     0     0      0
## [2,]           1     1     0      0
## [3,]           1     0     1      0
```

3.1 Alternatives

An alternative is to do:

```
c1 <- esticon(tooth1, L)
c1
```

```
##      estimate std.error statistic p.value  beta0 df
## [1,]   12.455     0.988    12.603   0.000  0.000 56
## [2,]   21.585     0.988    21.841   0.000  0.000 56
## [3,]   27.950     0.988    28.281   0.000  0.000 56
```

Notice: `esticon` has been in the **doBy** package for many years; `linest` is a newer addition; `esticon` is not actively maintained but remains in **doBy** for historical reasons. Yet another alternative in this case is to generate a new data frame and then invoke predict (but this approach is not generally applicable, see later):

```
nd <- data.frame(dose=c('0.5', '1', '2'), supp='OJ')
nd
```

```
##   dose supp
## 1  0.5   OJ
## 2    1   OJ
## 3    2   OJ
```

```
predict(tooth1, newdata=nd)
```

```
##     1     2     3
## 12.45 21.58 27.95
```

4 Least-squares means (LS–means)

A related question could be: What is the estimated length for each dose if we ignore the source of vitamin C (i.e. whether it is OJ or VC). One approach would be to fit a model in which source does not appear:

```
tooth0 <- update(tooth1, . ~ . - supp)
L0 <- LE_matrix(tooth0, effect="dose")
L0
```

```
##      (Intercept) dose1 dose2
## [1,]           1     0     0
## [2,]           1     1     0
## [3,]           1     0     1
```

```
linest(tooth0, L=L0)
```

```
## Coefficients:
##     estimate std.error statistic     df p.value
## [1,]   10.605     0.949    11.180 57.000       0
## [2,]   19.735     0.949    20.805 57.000       0
## [3,]   26.100     0.949    27.515 57.000       0
```

An alternative would be to stick to the original model but compute the estimate for an "average vitamin C source". That would correspond to giving weight $1/2$ to each of the two vitamin C source parameters. However, as one of the parameters is already set to zero to obtain identifiability, we obtain the LE–matrix L as

```
L1 <- matrix(c(1, 0, 0, 0.5,
               1, 1, 0, 0.5,
               1, 0, 1, 0.5), nrow=3, byrow=T)
linest(tooth1, L=L1)
```

```
## Coefficients:
##     estimate std.error statistic     df p.value
## [1,]   10.605     0.856    12.391 56.000       0
## [2,]   19.735     0.856    23.058 56.000       0
## [3,]   26.100     0.856    30.495 56.000       0
```

Such a particular linear estimate is sometimes called a least-squares mean or an LSmean or a marginal mean. Notice that the parameter estimates under the two approaches are identical. This is is because data is balanced: There are 10 observations per supplementation type. Had data not been balanced, the estimates would in general have been different.

Notice: One may generate L automatically with

```
L1 <- LE_matrix(tooth1, effect="dose")
L1
```

```
##      (Intercept) dose1 dose2 suppVC
## [1,]           1     0     0    0.5
## [2,]           1     1     0    0.5
## [3,]           1     0     1    0.5
```

Notice: One may obtain the LSmean directly as:

```
LSmeans(tooth1, effect="dose")

## Coefficients:
##        estimate std.error statistic      df p.value
## [1,]    10.605     0.856    12.391 56.000       0
## [2,]    19.735     0.856    23.058 56.000       0
## [3,]    26.100     0.856    30.495 56.000       0
```

which is the same as

```
L <- LE_matrix(tooth1, effect="dose")
le <- linest(tooth1, L=L)
coef(le)
```

For a model with interactions, the LSmeans are

```
LSmeans(tooth2, effect="dose")

## Coefficients:
##        estimate std.error statistic      df p.value
## [1,]    10.605     0.812    13.060 54.000       0
## [2,]    19.735     0.812    24.304 54.000       0
## [3,]    26.100     0.812    32.143 54.000       0
```

In this case, the LE–matrix is

```
L <- LE_matrix(tooth2, effect="dose")
t(L)

##               [,1] [,2] [,3]
## (Intercept)    1.0  1.0  1.0
## dose1          0.0  1.0  0.0
## dose2          0.0  0.0  1.0
## suppVC         0.5  0.5  0.5
## dose1:suppVC   0.0  0.5  0.0
## dose2:suppVC   0.0  0.0  0.5
```

5 Using the `at=` argument

```
library(ggplot2)
ChickWeight$Diet <- factor(ChickWeight$Diet)
qplot(Time, weight, data=ChickWeight, colour=Chick, facets=~Diet,
      geom=c("point","line"))
```

Consider random regression model:

Figure 3: ChickWeight data.

```
library(lme4)

## Loading required package:  Matrix

chick <- lmer(weight ~ Time * Diet + (0 + Time | Chick),
          data=ChickWeight)
coef(summary(chick))

##             Estimate Std. Error t value
## (Intercept)   33.218     1.7697 18.7701
## Time           6.339     0.6103 10.3855
## Diet2         -4.585     3.0047 -1.5258
## Diet3        -14.968     3.0047 -4.9815
## Diet4         -1.454     3.0177 -0.4818
## Time:Diet2     2.271     1.0367  2.1902
## Time:Diet3     5.084     1.0367  4.9043
## Time:Diet4     3.217     1.0377  3.1004
```

The LE–matrix for `Diet` becomes:

```
L <- LE_matrix(chick, effect="Diet")
t(L)

##               [,1]  [,2]  [,3]  [,4]
## (Intercept)   1.00  1.00  1.00  1.00
## Time         10.72 10.72 10.72 10.72
## Diet2         0.00  1.00  0.00  0.00
## Diet3         0.00  0.00  1.00  0.00
## Diet4         0.00  0.00  0.00  1.00
## Time:Diet2    0.00 10.72  0.00  0.00
## Time:Diet3    0.00  0.00 10.72  0.00
```

8

```
## Time:Diet4   0.00  0.00  0.00 10.72
```

The value of `Time` is by default taken to be the average of that variable. Hence the `LSmeans` is the predicted weight for each diet at that specific point of time. We can consider other points of time with

```
K1 <- LE_matrix(chick, effect="Diet", at=list(Time=1))
t(K1)
```

```
##              [,1] [,2] [,3] [,4]
## (Intercept)    1    1    1    1
## Time           1    1    1    1
## Diet2          0    1    0    0
## Diet3          0    0    1    0
## Diet4          0    0    0    1
## Time:Diet2     0    1    0    0
## Time:Diet3     0    0    1    0
## Time:Diet4     0    0    0    1
```

The `LSmeans` for the intercepts is the predictions at `Time=0`. The `LSmeans` for the slopes becomes

```
K0 <- LE_matrix(chick, effect="Diet", at=list(Time=0))
t(K1 - K0)
```

```
##              [,1] [,2] [,3] [,4]
## (Intercept)    0    0    0    0
## Time           1    1    1    1
## Diet2          0    0    0    0
## Diet3          0    0    0    0
## Diet4          0    0    0    0
## Time:Diet2     0    1    0    0
## Time:Diet3     0    0    1    0
## Time:Diet4     0    0    0    1
```

```
linest(chick, L=K1-K0)
```

```
## Coefficients:
##       estimate std.error statistic     df p.value
## [1,]    6.339     0.610   10.383 49.855       0
## [2,]    8.609     0.838   10.273 48.282       0
## [3,]   11.423     0.838   13.631 48.282       0
## [4,]    9.556     0.839   11.386 48.565       0
```

We can cook up our own function for comparing trends:

```
LSmeans_trend <- function(object, effect, trend){
    L <- LE_matrix(object, effect=effect, at=as.list(setNames(1, trend))) -
        LE_matrix(object, effect=effect, at=as.list(setNames(0, trend)))
    linest(object, L=L)
}
LSmeans_trend(chick, effect="Diet", trend="Time")
```

```
## Coefficients:
##      estimate std.error statistic     df p.value
## [1,]    6.339     0.610    10.383 49.855       0
## [2,]    8.609     0.838    10.273 48.282       0
## [3,]   11.423     0.838    13.631 48.282       0
## [4,]    9.556     0.839    11.386 48.565       0
```

6 Using (transformed) covariates

Consider the following subset of the `CO2` dataset:

```
data(CO2)
CO2 <- transform(CO2, Treat=Treatment, Treatment=NULL)
levels(CO2$Treat) <- c("nchil","chil")
levels(CO2$Type) <- c("Que","Mis")
ftable(xtabs( ~ Plant + Type + Treat, data=CO2), col.vars=2:3)
```

```
##         Type   Que       Mis
##        Treat nchil chil nchil chil
## Plant
## Qn1            7    0     0    0
## Qn2            7    0     0    0
## Qn3            7    0     0    0
## Qc1            0    7     0    0
## Qc3            0    7     0    0
## Qc2            0    7     0    0
## Mn3            0    0     7    0
## Mn2            0    0     7    0
## Mn1            0    0     7    0
## Mc2            0    0     0    7
## Mc3            0    0     0    7
## Mc1            0    0     0    7
```

```
qplot(x=log(conc), y=uptake, data=CO2, color=Treat, facets=~Type)
```

Below, the covariate `conc` is fixed at the average value:

```
co2.lm1 <- lm(uptake ~ conc + Type + Treat, data=CO2)
LSmeans(co2.lm1, effect="Treat")
```

```
## Coefficients:
##      estimate std.error statistic     df p.value
## [1,]   30.643     0.956    32.066 80.000       0
## [2,]   23.783     0.956    24.888 80.000       0
```

If we use `log(conc)` instead we will get an error when calculating LS–means:

Figure 4: CO2 data

```
co2.lm <- lm(uptake ~ log(conc) + Type + Treat, data=CO2)
LSmeans(co2.lm, effect="Treat")
```

In this case one can do

```
co2.lm2 <- lm(uptake ~ log.conc + Type + Treat,
              data=transform(CO2, log.conc=log(conc)))
LSmeans(co2.lm2, effect="Treat")

## Coefficients:
##      estimate std.error statistic      df p.value
## [1,]   30.643     0.761    40.261 80.000       0
## [2,]   23.783     0.761    31.248 80.000       0
```

This also highlights what is computed: The average of the log of `conc`; not the log of the average of `conc`.

In a similar spirit consider

```
co2.lm3 <- lm(uptake ~ conc + I(conc^2) + Type + Treat, data=CO2)
LSmeans(co2.lm3, effect="Treat")

## Coefficients:
##      estimate std.error statistic      df p.value
## [1,]   34.543     0.982    35.191 79.000       0
## [2,]   27.683     0.982    28.202 79.000       0
```

Above `I(conc^2)` is the average of the squared values of `conc`; not the square of the average of `conc`, cfr. the following.

```
co2.lm4 <- lm(uptake ~ conc + conc2 + Type + Treat, data=
              transform(CO2, conc2=conc^2))
LSmeans(co2.lm4, effect="Treat")

## Coefficients:
##      estimate std.error statistic      df p.value
## [1,]   30.643     0.776    39.465 79.000       0
## [2,]   23.783     0.776    30.630 79.000       0
```

If we want to evaluate the LS–means at `conc=10` then we can do:

```
LSmeans(co2.lm4, effect="Treat", at=list(conc=10, conc2=100))

## Coefficients:
##      estimate std.error statistic     df p.value
## [1,]    14.74      1.70      8.66 79.00       0
## [2,]     7.88      1.70      4.63 79.00       0
```

7 Alternative models

7.1 Generalized linear models

We can calculate LS–means for e.g. a Poisson or a gamma model. Default is that the calculation
is calculated on the scale of the linear predictor. However, if we think of LS–means as a prediction
on the linear scale one may argue that it can also make sense to transform this prediction to the
response scale:

```
tooth.gam <- glm(len ~ dose + supp, family=Gamma, data=ToothGrowth)
LSmeans(tooth.gam, effect="dose", type="link")

## Coefficients:
##      estimate std.error statistic p.value
## [1,]  0.09453   0.00579  16.33340       0
## [2,]  0.05111   0.00312  16.39673       0
## [3,]  0.03889   0.00238  16.36460       0

LSmeans(tooth.gam, effect="dose", type="response")

## Coefficients:
##      estimate std.error statistic p.value
## [1,]  0.09453   0.00579  16.33340       0
## [2,]  0.05111   0.00312  16.39673       0
## [3,]  0.03889   0.00238  16.36460       0
```

7.2 Linear mixed effects model

For the sake of illustration we treat `supp` as a random effect:

```
library(lme4)
tooth.mm <- lmer( len ~ dose  + (1|supp), data=ToothGrowth)
LSmeans(tooth1, effect="dose")


## Coefficients:
##      estimate std.error statistic      df p.value
## [1,]   10.605     0.856    12.391 56.000       0
## [2,]   19.735     0.856    23.058 56.000       0
## [3,]   26.100     0.856    30.495 56.000       0


LSmeans(tooth.mm, effect="dose")


## Coefficients:
##      estimate std.error statistic    df p.value
## [1,]    10.61      1.98      5.36  1.31    0.08
## [2,]    19.74      1.98      9.98  1.31    0.03
## [3,]    26.10      1.98     13.20  1.31    0.02
```

Notice here that the estimates themselves identical to those of a linear model (that is not generally the case, but it is so here because data is balanced). In general the estimates are will be very similar but the standard errors are much larger under the mixed model. This comes from that there that `supp` is treated as a random effect.

```
VarCorr(tooth.mm)


##  Groups   Name        Std.Dev.
##  supp     (Intercept) 2.52
##  Residual             3.83
```

Notice that the degrees of freedom by default are adjusted using a Kenward–Roger approximation (provided that **pbkrtest** is installed). Unadjusted degrees of freedom are obtained by setting `adjust.df=FALSE`.

7.3   Generalized estimating equations

Lastly, for gee-type "models" we get

```
library(geepack)
tooth.gee <- geeglm(len ~ dose, id=supp, family=Gamma, data=ToothGrowth)
LSmeans(tooth.gee, effect="dose")


## Coefficients:
##      estimate std.error statistic p.value
## [1,] 9.43e-02  1.65e-02  5.71e+00       0
## [2,] 5.07e-02  5.38e-03  9.41e+00       0
## [3,] 3.83e-02  4.15e-05  9.23e+02       0


LSmeans(tooth.gee, effect="dose", type="response")


## Coefficients:
```

```
##      estimate std.error statistic p.value
## [1,] 9.43e-02  1.65e-02  5.71e+00        0
## [2,] 5.07e-02  5.38e-03  9.41e+00        0
## [3,] 3.83e-02  4.15e-05  9.23e+02        0
```

8 Miscellaneous

8.1 Example: Non–estimable linear functions

```
## Make balanced dataset
dat.bal <- expand.grid(list(AA=factor(1:2), BB=factor(1:3), CC=factor(1:3)))
dat.bal$y <- rnorm(nrow(dat.bal))

## Make unbalanced dataset:  'BB' is nested within 'CC' so BB=1
## is only found when CC=1 and BB=2,3 are found in each CC=2,3,4
dat.nst <- dat.bal
dat.nst$CC <-factor(c(1,1,2,2,2,2,1,1,3,3,3,3,1,1,4,4,4,4))
```

```
dat.nst
```

```
##     AA BB CC       y
## 1    1  1  1 -0.2176
## 2    2  1  1  1.3518
## 3    1  2  2  0.2363
## 4    2  2  2  2.8840
## 5    1  3  2 -0.2189
## 6    2  3  2  0.7217
## 7    1  1  1 -0.7102
## 8    2  1  1  0.7570
## 9    1  2  3  1.6901
## 10   2  2  3 -0.6459
## 11   1  3  3 -1.9504
## 12   2  3  3 -1.7034
## 13   1  1  1 -0.7973
## 14   2  1  1 -0.2924
## 15   1  2  4  0.7295
## 16   2  2  4 -1.9058
## 17   1  3  4 -0.5614
## 18   2  3  4 -1.3506
```

Consider this simulated dataset:

```
head(dat.nst, 4)
```

```
##   AA BB CC       y
## 1  1  1  1 -0.2176
## 2  2  1  1  1.3518
## 3  1  2  2  0.2363
## 4  2  2  2  2.8840
```

```
ftable(xtabs( ~ AA + BB + CC, data=dat.nst), row.vars="AA")
```

```
##    BB 1        2        3
##    CC 1 2 3 4 1 2 3 4 1 2 3 4
## AA
## 1     3 0 0 0 0 1 1 1 0 1 1 1
## 2     3 0 0 0 0 1 1 1 0 1 1 1
```

Data is highly "unbalanced": Whenever `BB=1` then `CC` is always `1`; whenever `BB` is not `1` then `CC` is never `1`. We have

```
mod.nst  <- lm(y ~ AA + BB : CC, data=dat.nst)
coef(summary(mod.nst))
```

```
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -1.0458     0.8825 -1.1851   0.2634
## AA2           0.1796     0.5581  0.3218   0.7542
## BB1:CC1       0.9712     0.9667  1.0047   0.3387
## BB2:CC2       2.5161     1.1839  2.1252   0.0595
## BB3:CC2       1.2074     1.1839  1.0198   0.3319
## BB2:CC3       1.4781     1.1839  1.2485   0.2403
## BB3:CC3      -0.8710     1.1839 -0.7356   0.4789
## BB2:CC4       0.3678     1.1839  0.3106   0.7624
```

In this case some of the `LSmeans` values are not estimable; for example:

```
lsm.BC <- LSmeans(mod.nst, effect=c("BB", "CC"))
lsm.BC
```

```
## Coefficients:
##       estimate std.error statistic      df p.value
##  [1,]   0.0152    0.4833    0.0315 10.0000    0.98
##  [2,]       NA        NA        NA      NA      NA
##  [3,]       NA        NA        NA      NA      NA
##  [4,]       NA        NA        NA      NA      NA
##  [5,]   1.5601    0.8372    1.8636 10.0000    0.09
##  [6,]   0.2514    0.8372    0.3003 10.0000    0.77
##  [7,]       NA        NA        NA      NA      NA
##  [8,]   0.5221    0.8372    0.6237 10.0000    0.55
##  [9,]  -1.8269    0.8372   -2.1823 10.0000    0.05
## [10,]       NA        NA        NA      NA      NA
## [11,]  -0.5882    0.8372   -0.7026 10.0000    0.50
## [12,]  -0.9560    0.8372   -1.1419 10.0000    0.28
```

```
lsm.BC2 <- LSmeans(mod.nst, effect="BB", at=list(CC=2))
lsm.BC2
```

```
## Coefficients:
##       estimate std.error statistic      df p.value
## [1,]       NA        NA        NA      NA      NA
## [2,]    1.560     0.837     1.864 10.000    0.09
## [3,]    0.251     0.837     0.300 10.000    0.77
```

We describe the situation in Section 8.2 where we focus on `lsm.BC2`.

8.2 Handling non–estimability

The model matrix for the model in Section 8.1 does not have full column rank and therefore not all values are calculated by `LSmeans()`.

```
X <- model.matrix( mod.nst )
Matrix::rankMatrix(X)


## [1] 8
## attr(,"method")
## [1] "tolNorm2"
## attr(,"useGrad")
## [1] FALSE
## attr(,"tol")
## [1] 3.997e-15


dim(X)


## [1] 18 14


as(X, "Matrix")


## 18 x 14 sparse Matrix of class "dgCMatrix"


##    [[ suppressing 14 column names '(Intercept)', 'AA2', 'BB1:CC1' ...  ]]


##
## 1   1 . 1 . . . . . . . . . . .
## 2   1 1 1 . . . . . . . . . . .
## 3   1 . . . . . 1 . . . . . . .
## 4   1 1 . . . . 1 . . . . . . .
## 5   1 . . . . . . 1 . . . . . .
## 6   1 1 . . . . . 1 . . . . . .
## 7   1 . 1 . . . . . . . . . . .
## 8   1 1 1 . . . . . . . . . . .
## 9   1 . . . . . . . . 1 . . . .
## 10  1 1 . . . . . . . 1 . . . .
## 11  1 . . . . . . . . . 1 . . .
## 12  1 1 . . . . . . . . 1 . . .
## 13  1 . 1 . . . . . . . . . . .
## 14  1 1 1 . . . . . . . . . . .
## 15  1 . . . . . . . . . . . 1 .
## 16  1 1 . . . . . . . . . . 1 .
## 17  1 . . . . . . . . . . . . 1
## 18  1 1 . . . . . . . . . . . 1
```

We consider a model, i.e. an n dimensional random vector $y = (y_i)$ for which $\mathbb{E}(y) = \mu = X\beta$ and $\mathbb{C}\mathrm{ov}(y) = V$ where X does not have full column rank We are interested in linear functions of β,

say
$$c = l^\top \beta = \sum_j l_j \beta_j.$$

```
L <- LE_matrix(mod.nst, effect="BB", at=list(CC=2))
t(L)
```

```
##                [,1] [,2] [,3]
## (Intercept)  1.0  1.0  1.0
## AA2          0.5  0.5  0.5
## BB1:CC1      0.0  0.0  0.0
## BB2:CC1      0.0  0.0  0.0
## BB3:CC1      0.0  0.0  0.0
## BB1:CC2      1.0  0.0  0.0
## BB2:CC2      0.0  1.0  0.0
## BB3:CC2      0.0  0.0  1.0
## BB1:CC3      0.0  0.0  0.0
## BB2:CC3      0.0  0.0  0.0
## BB3:CC3      0.0  0.0  0.0
## BB1:CC4      0.0  0.0  0.0
## BB2:CC4      0.0  0.0  0.0
## BB3:CC4      0.0  0.0  0.0
```

```
linest(mod.nst, L=L)
```

```
## Coefficients:
##      estimate std.error statistic      df p.value
## [1,]       NA        NA        NA      NA      NA
## [2,]    1.560     0.837     1.864  10.000    0.09
## [3,]    0.251     0.837     0.300  10.000    0.77
```

A least squares estimate of β is
$$\hat{\beta} = GX^\top y$$

where G is a generalized inverse of $X^\top X$. Since the generalized inverse is not unique then neither is the estimate $\hat{\beta}$. Hence $\hat{c} = l^\top \hat{\beta}$ is in general not unique.

One least squares estimate of β and one corresponding linear estimate $L\hat{\beta}$ is:

```
XtXinv <- MASS::ginv(t(X)%*%X)
bhat <- as.numeric(XtXinv %*% t(X) %*% dat.nst$y)
zapsmall(bhat)
```

```
##  [1] -0.2063  0.1796  0.1318  0.0000  0.0000  0.0000  1.6767  0.3679  0.0000  0.6387
## [11] -1.7104  0.0000 -0.4716 -0.8394
```

```
L %*% bhat
```

```
##          [,1]
## [1,] -0.1166
## [2,]  1.5601
## [3,]  0.2514
```

For some values of l (i.e. for some rows of L) the estimate $\hat{c} = l^\top\beta$ is unique (i.e. it does not depend on the choice of generalized inverse). Such linear functions are said to be estimable and can be described as follows:

All we specify with $\mu = X\beta$ is that μ is a vector in the column space $C(X)$ of X. We can only learn about β through $X\beta$ so the only thing we can say something about is linear combinations $\rho^\top X\beta$. Hence we can only say something about $l^\top\beta$ if there exists ρ such that

$$l^\top\beta = \rho^\top X\beta,$$

i.e., if $l = X^\top\rho$ for some ρ, which is if l is in the column space $C(X^\top)$ of $X^\top$. This is the same as saying that l must be perpendicular to all vectors in the null space $N(X)$ of X. To check this, we find a basis B for $N(X)$. This can be done in many ways, for example via a singular value decomposition of X, i.e.

$$X = UDV^\top$$

A basis for $N(X)$ is given by those columns of V that corresponds to zeros on the diagonal of D.

```
S <- svd(X)
B <- S$v[, S$d < 1e-10, drop=FALSE ];
head(B) ## Basis for N(X)
```

```
##              [,1]       [,2]       [,3]       [,4]       [,5]       [,6]
## [1,] -3.245e-01  7.275e-02  2.114e-02  7.406e-02  9.204e-02  0.000e+00
## [2,]  1.110e-16 -2.776e-17  1.388e-17 -2.776e-17 -2.776e-17 -2.105e-48
## [3,]  3.245e-01 -7.275e-02 -2.114e-02 -7.406e-02 -9.204e-02  2.784e-48
## [4,] -3.655e-01 -1.092e-01 -1.102e-01 -5.458e-01 -7.378e-01  1.181e-18
## [5,]  8.250e-02  5.599e-01 -6.591e-01  3.858e-01 -3.107e-01 -4.748e-17
## [6,] -1.228e-01 -6.891e-01 -6.758e-01  5.269e-02  2.247e-01 -3.182e-17
```

From

```
rowSums(L %*% B)
```

```
## [1] -1.275e+00  2.658e-15 -1.256e-15
```

we conclude that the first row of L is not perpendicular to all vectors in the null space $N(X)$ whereas the two last rows of L are. Hence these two linear estimates are estimable; their value does not depend on the choice of generalized inverse:

```
lsm.BC2
```

```
## Coefficients:
##       estimate std.error statistic      df p.value
## [1,]       NA        NA        NA      NA      NA
## [2,]    1.560     0.837     1.864  10.000    0.09
## [3,]    0.251     0.837     0.300  10.000    0.77
```

8.3 Pairwise comparisons

We will just mention that for certain other linear estimates, the matrix L can be generated automatically using `glht()` from the **multcomp** package. For example, pairwise comparisons of all levels of `dose` can be obtained with

```
library("multcomp")
```

```
## Loading required package:  mvtnorm
## Loading required package:  survival
## Loading required package:  TH.data
## Loading required package:  MASS
##
## Attaching package:  'TH.data'
## The following object is masked from 'package:MASS':
##
##    geyser
```

```
g1 <- glht(tooth1, mcp(dose="Tukey"))
summary( g1 )
```

```
##
##   Simultaneous Tests for General Linear Hypotheses
##
## Multiple Comparisons of Means: Tukey Contrasts
##
##
## Fit: lm(formula = len ~ dose + supp, data = ToothGrowth)
##
## Linear Hypotheses:
##             Estimate Std. Error t value Pr(>|t|)
## 1 - 0.5 == 0     9.13       1.21    7.54   <1e-05 ***
## 2 - 0.5 == 0    15.50       1.21   12.80   <1e-05 ***
## 2 - 1 == 0       6.37       1.21    5.26   <1e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## (Adjusted p values reported -- single-step method)
```

The L matrix is

```
L <- g1$linfct
L
```

```
##          (Intercept) dose1 dose2 suppVC
## 1 - 0.5            0     1     0      0
## 2 - 0.5            0     0     1      0
## 2 - 1             0    -1     1      0
## attr(,"type")
## [1] "Tukey"
```

and this matrix can also be supplied to `glht`

```
glht(tooth1, linfct=L)
```